

Towards Better 3D Knowledge Transfer via Masked Image Modeling for Multi-view 3D Understanding

Jihao Liu^{1,2} Tai Wang^{1,3} Boxiao Liu² Qihang Zhang¹ Yu Liu² Hongsheng Li^{1,4}
¹CUHK MMLab ²SenseTime Research
³Shanghai AI Laboratory ⁴CPII under InnoHK

Abstract

Multi-view camera-based 3D detection is a challenging problem in computer vision. Recent works leverage a pre-trained LiDAR detection model to transfer knowledge to a camera-based student network. However, we argue that there is a major domain gap between the LiDAR BEV features and the camera-based BEV features, as they have different characteristics and are derived from different sources. In this paper, we propose Geometry Enhanced Masked Image Modeling (GeoMIM) to transfer the knowledge of the LiDAR model in a pretrain-finetune paradigm for improving the multi-view camera-based 3D detection. GeoMIM is a multi-camera vision transformer with Cross-View Attention (CVA) blocks that uses LiDAR BEV features encoded by the pre-trained BEV model as learning targets. During pretraining, GeoMIM’s decoder has a semantic branch completing dense perspective-view features and the other geometry branch reconstructing dense perspective-view depth maps. The depth branch is designed to be camera-aware by inputting the camera’s parameters for better transfer capability. Extensive results demonstrate that GeoMIM outperforms existing methods on nuScenes benchmark, achieving state-of-the-art performance for camera-based 3D object detection and 3D segmentation.

1. Introduction

Multi-view camera-based 3D detection is an emerging critical problem in computer vision [20, 44, 45, 19, 28, 29, 36, 26, 31, 32]. To improve the detection performance, recent works [9, 27, 21] often choose to use a pretrained LiDAR model as the teacher and transfer its knowledge to a camera-based student network. Various techniques, such as LIGA-Stereo [13], CMKD [17], and BEVDistill [8], have been proposed to leverage the rich geometry information of the LiDAR model’s BEV (bird’s eye view) features.

Utilizing a pretrained LiDAR model to provide auxiliary supervision has become a widely adopted design that can enhance the performance of camera-based models. However, we contend that this design is not optimal due to a

Pretrain	Supervision	Finetune	Finetune + LiDAR BEV
SL [33]	Classes	40.6	41.7
SSL [30]	RGB Pixels	44.3	43.9
GeoMIM	BEV Feature	47.2	45.4

Table 1: The effects of LiDAR BEV feature distillation on ImageNet-pretrained (SL), self-supervised (SSL), and our GeoMIM pretraining-finetuning settings for BEVDet in nuScenes 3D detection. Naively distilling LiDAR BEV features in finetuning introduces domain gaps and harms the performance when the pretrained model is powerful enough.

significant domain gap between the BEV features of the LiDAR model and those of the camera-based model. This domain gap arises from the 3D and sparse characteristics of LiDAR point clouds compared to the dense 2D images captured by the camera. Additionally, the LiDAR model’s BEV features are grounded in ground truth depth, while those of the camera-based model are typically inferred from 2D images, a problem that is often ill-posed. We empirically demonstrate their domain gap with a pilot study as shown in Tab. 1. We find that utilizing a LiDAR teacher to provide auxiliary supervision can indeed improve an ImageNet-pretrained [33] camera-based model, but is unable to improve a stronger camera-based model initialized by recent powerful self-supervised pretraining. In other words, directly utilizing the pretrained LiDAR model to distill the final camera-based model might not be an optimal design and does not necessarily lead to performance gain.

To better take advantage of the LiDAR model, in this paper, we propose *Geometry Enhanced Masked Image Modeling (GeoMIM)* to transfer the knowledge of the LiDAR model in a pretrain-finetune paradigm for improving the multi-view camera-based 3D detection. It is built upon a multi-camera vision transformer with *Cross-View Attention (CVA)* blocks and enables perspective-view (PV) representation pretraining via BEV feature reconstruction from masked images. Specifically, during pretraining, we partition the training images into patches and feed a portion of them into the encoder following Masked Autoencoder [14]. Our GeoMIM decoder then uses these encoded visible to-

kens to reconstruct the pretrained LiDAR model’s BEV feature in the BEV space instead of commonly used RGB pixels [47, 14, 30] or depth points [3] as in existing MAE frameworks. To achieve this PV to BEV reconstruction, we first devise two branches to *decouple* the semantic and geometric parts, with one branch completing dense PV features and the other reconstructing the depth map. The dense PV features can then be projected into the BEV space with the depth distribution following Lift-Splat-Shoot (LSS) [37]. We further equip the two branches with the proposed CVA blocks in their intermediate layers to allow each patch to attend to tokens in other views. It enhances the decoder’s capability of joint multi-view inference which is especially critical for BEV feature reconstruction. Finally, the depth branch is designed to be *camera-aware* with the additional encoding of cameras’ parameters as input, making the pretrained GeoMIM better adapt to downstream tasks with different cameras.

To demonstrate the effectiveness of GeoMIM, we fine-tune the pretrained backbone to conduct multi-view camera-based 3D detection and 3D segmentation on the nuScenes [7] dataset. We achieve state-of-the-art results of 64.4 NDS (NuScenes Detection Score) and 70.5 mIoU (mean intersection over union) for 3D detection and segmentation on the NuScenes `test` set, which are 2.5% and 1.1% better than previously reported best results [36, 22]. Additionally, we verify that the backbone pretrained on nuScenes dataset can be successfully transferred to Waymo Open dataset [40], improving the mAP (mean average precision) of the ImageNet-initialized 3D detector by 6.9%.

2. Related Works

Masked Image Modeling Inspired by BERT [11] for Masked Language Modeling, Masked Image Modeling (MIM) becomes a popular pretext task for visual representation learning [6, 14, 2, 46, 1, 4, 51, 3, 49]. MIM aims to reconstruct the masked tokens from a corrupted input. SimMIM [47] points out that raw pixel values of the randomly masked patches are a good reconstruction target and a lightweight prediction head is sufficient for pretraining. Different from SimMIM, MAE [14] only takes the visible patches as the input of the encoder. Mask tokens are added in the middle of the encoder and the decoder. BEiT [6] utilizes a pretrained discrete VAE (dVAE) [39, 38] as the tokenizer. PeCo [12] proposed to apply perceptual similarity loss on the training of dVAE can drive the tokenizer to generate better semantic visual tokens, which helps pretraining. In contrast to those works, our GeoMIM utilizes a geometry-rich LiDAR model and transfers its knowledge via MIM pretraining, aiming to improve the multi-view camera-based 3D models.

Multi-view camera-based 3D detection The field of camera-based 3D object detection has seen significant progress in recent years [44, 45, 29, 20, 28, 26, 36, 50].

FCOS3D [44] proposed a fully convolutional single-stage detector for monocular 3D object detection. DETR3D [45] extends the DETR framework to the 3D domain, and proposes a framework for end-to-end 3D object detection. BEVFormer [29] combines BEV (bird’s eye view) representation and transformer networks for 3D object detection. BEVDepth [28] focuses on accurately estimating the depth of objects in the BEV representation. Additionally, considering the promising performance of the LiDAR-based detectors, there are several papers that use a pretrained LiDAR detector for knowledge distillation [16]. LIGA-Stereo [13] proposes to mimic the LiDAR BEV features for training a camera-based detector. UVTR [27] represents different modalities in a unified manner and supports knowledge transfer with the voxel representations. More recent BEVDistill [8] and CMKD [17] not only use the LiDAR BEV features for knowledge distillation but also transfer the teacher’s knowledge through sparse instance distillation and response-based distillation respectively. In comparison, we utilized the pretrained LiDAR model in a pretraining-finetuning paradigm to avoid the LiDAR-camera BEV domain gap.

3. Method

Employing a pretrained LiDAR-based detection model to provide auxiliary learning guidance to train camera-based 3D understanding models has shown promising results in recent years [13, 8, 9, 27, 17]. However, because of the domain gap between the LiDAR and camera modalities, we observe that when a camera-based model is already strong, directly supervising it with the LiDAR teacher fails to improve the camera-based model as shown in Tab. 1.

To address this problem, we propose GeoMIM to better transfer the LiDAR model’s knowledge to the camera-based model in a pretrain-finetune paradigm. GeoMIM pretrains a multi-view camera-based model via Masked Image Modeling (MIM) [47]. Unlike existing 2D MAE works [14, 30], we project the semantic features to the BEV (bird’s eye view) space and use the LiDAR BEV features in the 3D space as the reconstruction targets for pretraining. The pretrained LiDAR model is only used in the pretraining stage, and is discarded in the finetuning stage to avoid introducing the LiDAR-camera BEV domain gap. We illustrate the proposed GeoMIM in Fig. 1.

Masking and Encoder Given the multi-view input images $X = \{x_i \in \mathbb{R}^{3 \times H \times W}, i = 1, 2, \dots, N\}$ where N , H , W are the number of views, image height, and width, we randomly mask a proportion of input image patches (tokens) and use a Swin Transformer [33] as the encoder to encode the visible tokens. The encoded representations, $F^v \in \mathbb{R}^{N \times C \times L}$ where C and L denote the number of dimensions and the number of visible tokens, are then filled with a shared mask token $[M] \in \mathbb{R}^C$ at the masked locations and further processed by the decoder for reconstruction.

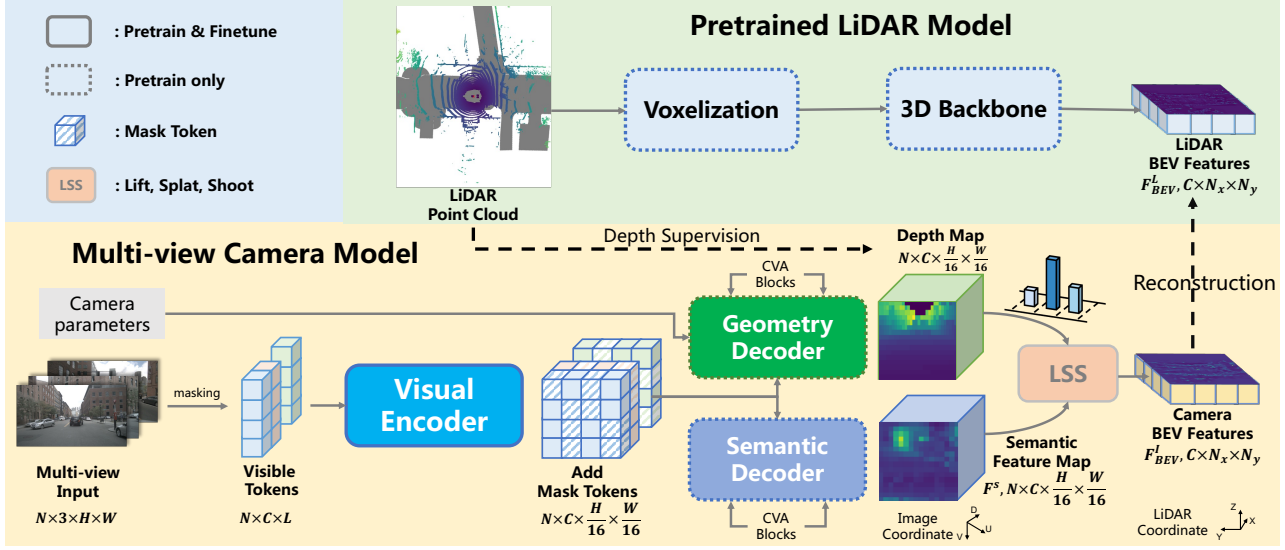


Figure 1: Overview of GeoMIM. For pretraining, the multi-view images are randomly masked for a proportion of image tokens, and only the visible tokens are processed by the encoder. Right before decoding, the token embeddings are filled with mask tokens for separately decoding dense camera-view semantic features and depth maps, which are then projected to BEV space for reconstructing the LiDAR BEV features. After pretraining, only the encoder is finetuned on downstream tasks.




Figure 2: Cross-view attention block. We partition the multi-view inputs into multiple groups according to their row indices, and perform self-attention within each group.

GeoMIM Decoder To transfer the rich geometry knowledge of a pretrained LiDAR detector to our camera-based model, we jointly project the multi-view semantic features according to their estimated depth maps to the BEV space and use the same scene’s LiDAR BEV features as the reconstruction targets. Specifically, our GeoMIM uses two *decoupled* decoders, each of which consists of 8 Transformer [41] blocks. The semantic decoder D_{sem} reconstructs the dense camera-view semantic features $F^s \in \mathbb{R}^{N \times C \times \frac{H}{16} \times \frac{W}{16}}$ of the N camera views and the other geometry decoder D_{geo} predicts dense camera-view depth maps $D \in \mathbb{R}^{N \times B \times \frac{H}{16} \times \frac{W}{16}}$ of the N camera views, where B denotes the number of depth bins. The depth map and semantic feature can be expressed as

$$D = D_{\text{geo}}(F^v, [M]), \quad F^s = D_{\text{sem}}(F^v, [M]). \quad (1)$$

We can then obtain the camera BEV features F_{BEV}^I by jointly projecting the multi-view semantic features to the BEV space with the Lift-Splat-Shoot (LSS) [37] operation

according to the predicted dense depth maps,

$$F_{\text{BEV}}^I \in \mathbb{R}^{C \times N_x \times N_y} = \text{LSS}(F^s, D), \quad (2)$$

where N_x, N_y are the numbers of bins in the x and y axis of the BEV feature maps respectively. Empirically, the two decoders share the first half of the Transformer blocks for efficiency.

Unlike existing works that separately process the multi-view input images, we propose a novel *Cross-View Attention* (CVA) block to model the interaction across different views to better reconstruct the LiDAR BEV features from input images. Our intuition is that as the multi-view images are naturally overlapped, proper interaction across views is beneficial to align those images and better infer the LiDAR BEV features. Instead of explicitly using the epipolar lines to associate pixels across the multi-view images, we partition the camera-view tokens of the multiple views into groups according to their row indices and only allow the tokens belonging to the same row of the $\frac{1}{16}$ input resolution to interact with each other. The interaction is modeled by the self-attention operation [41]. Notably, our proposed CVA has linear computation complexity to the input image size and is therefore much more efficient compared to global self-attention. We illustrate the proposed CVA in Fig. 2. We use the CVA block as the 2th and 6th attention blocks of the decoder. Note that we do not add it to the backbone and no extra computation is introduced when finetuning the encoder.

Accurately reconstructing depth with the geometry decoder implicitly requires the decoder to infer the camera’s intrinsic parameters, which is difficult to generalize to an unseen dataset as the data may be collected with different

cameras. To achieve better transferability across different downstream tasks, we encode the camera’s intrinsic and extrinsic parameters using a linear projection layer and use the resulting features to scale the geometry decoder’s feature using the Squeeze-and-Excitation module [18]. Importantly, we do not require the camera’s information when finetuning on downstream tasks since only the decoder uses the camera’s information during pretraining. We demonstrate that the camera-aware depth reconstruction branch leads to better performance when finetuning on tasks that differ from the pretraining dataset.

Loss We use the mean squared error (MSE) loss between the projected camera BEV features and the pretrained LiDAR BEV features for pretraining,

$$\mathcal{L}_{rec} = \|(F_{BEV}^I - F_{BEV}^L)\|_2^2, \quad (3)$$

where $F_{BEV}^L \in \mathbb{R}^{C \times N_x \times N_y}$ denotes the pretrained LiDAR model’s BEV features. In addition, we incorporate a depth prediction task. Following prior arts [28], we use the ground truth discrete depth D_{GT} derived from the LiDAR point cloud and calculate the binary cross entropy (BCE) loss as the depth loss,

$$\mathcal{L}_{depth} = \text{BCE}(D, D_{GT}). \quad (4)$$

The overall loss can be expressed as

$$\mathcal{L} = \mathcal{L}_{rec} + \alpha \mathcal{L}_{depth}, \quad (5)$$

where α balances the two loss terms, which is set as 0.01 experimentally. Empirically, we observe that the depth loss can enhance the convergence speed, which is crucial for pretraining large models.

After pretraining, we discard the decoders and add a task-specific head on the top of the encoder for downstream tasks finetuning. During finetuning, we only utilize ground-truth supervision and abstain from utilizing the LiDAR model to avoid introducing the aforementioned domain gap.

Comparison with 2D MAE Compared to existing 2D MAE models [14, 47, 30], our proposed GeoMIM’s pretraining has two distinct characteristics: (1) We employ a geometry-rich LiDAR model and transfer its high-level knowledge in the BEV space via MIM pretraining, which can effectively enhance the geometry perception capability of the camera-based model. In contrast, the original MAE [14] reconstructs image pixels and could work well for 2D downstream perception tasks, but is found to be less effective for 3D perception. The reason is that the autonomous driving dataset, e.g., nuScenes [7], is much less diverse than MAE’s pretraining dataset ImageNet-1K [10]. As a result, employing image pixel reconstruction as the pretext task is hard to learn high-quality representations. (2) Contrary to MAE which only calculates the reconstruction loss in the masked tokens, we take all tokens into consideration in our loss. This is because the learning targets we use are from a different

modality and in a different geometric space. We can take full advantage of the LiDAR model by using all tokens to calculate the loss. For the masked locations, the objective is a prediction task while for the unmasked locations, it is similar to a distillation task.

4. Experiment Setups

To demonstrate the effectiveness of GeoMIM, we conduct experiments by pretraining Swin Transformer [33] backbones with GeoMIM and then finetuning it on various downstream tasks. These tasks include multi-view camera-based 3D detection on nuScenes [7] and Open Waymo [40] datasets, camera-based 3D semantic segmentation on nuScenes dataset, and 2D detection on nuImages dataset.

Dataset and Evaluation Metrics We use the large-scale nuScenes dataset for pretraining and finetuning, which contains 750, 150, and 150 scenes for training, validation, and testing, respectively. Each scene has 6 camera images and LiDAR point cloud covering 360°. Following the official evaluation metrics, we primarily report NuScenes Detection Score (NDS) and mean Average Precision (mAP) for comparison. We also report other five metrics, including ATE, ASE, AOE, AVE, and AAE, to measure translation, scale, orientation, velocity, and attribute errors, respectively, for a more detailed diagnosis.

We also evaluate the transferability of GeoMIM by finetuning the pretrained backbone on the Open Waymo and nuImages datasets. We report LET-3D-APL [23] and LET-3D-AP following the latest official guidelines for comparison. We report Mean Average Precision (mAP) of box and mask on nuImages dataset for 2D object detection and instance segmentation.

Pretraining We pretrain the Swin Transformer backbones on the training split of the nuScenes dataset with multi-view images as input. By default, we pretrain for 50 epochs with an input size of 256×704 . For ablation studies, we pretrain for 6 epochs unless otherwise specified. We use a pretrained TransFusion-L [5] LiDAR model to provide the reconstruction targets. We randomly mask the multi-view input images with a mask ratio of 50%. We use AdamW [34] optimizer with a learning rate of 2×10^{-4} and weight decay of 0.01. The learning rate is linearly warmed-up for 500 iterations and cosine decayed to 0. We apply the data augmentation strategy in BEVDet [20] to augment the input images and do not use augmentations for the LiDAR inputs. We utilize the Swin-Base and -Large backbones for pretraining, initializing the backbone with self-supervised ImageNet-pretraining [30].

Finetuning We keep the pretrained encoder, abandon the decoders, and adopt state-of-the-art frameworks for finetuning. We mainly evaluate the performance of the finetuned models on the 3D detection task on the nuScenes dataset. We also assess the transferability of GeoMIM on other down-

Framework	Pretrain	Backbone	Image Size	CBGS	mAP \uparrow	NDS \uparrow	mATE \downarrow	mASE \downarrow	mAOE \downarrow	mAVE \downarrow	mAAE \downarrow
DETR3D [45]	FCOS3D	R101-DCN	900 \times 1600	✓	0.349	0.434	0.716	0.268	0.379	0.842	0.200
BEVFormer [29]			900 \times 1600	✗	0.416	0.517	0.673	0.274	0.372	0.394	0.198
UVTR [27]			900 \times 1600	✗	0.379	0.483	0.731	0.267	0.350	0.510	0.200
PolarFormer [24]			900 \times 1600	✗	0.432	0.528	0.648	0.270	0.348	0.409	0.201
PETR [31]	ImageNet	R101	512 \times 1408	✓	0.357	0.421	0.710	0.270	0.490	0.885	0.224
PETRv2 [32]			640 \times 1600	✓	0.421	0.524	0.681	0.267	0.357	0.377	0.186
SOLOFusion [36]			512 \times 1408	✓	0.483	0.582	0.503	0.264	0.381	0.246	0.207
BEVDepth [28]	ImageNet	ConvNeXt-B	512 \times 1408	✓	0.462	0.558	-	-	-	-	-
BEVStereo [26]			512 \times 1408	✓	0.478	0.575	-	-	-	-	-
BEVDet4D [19]	ImageNet	Swin-B	640 \times 1600	✓	0.421	0.545	0.579	0.258	0.329	0.301	0.191
BEVDepth †			512 \times 1408	✓	0.466	0.555	0.531	0.264	0.489	0.293	0.200
BEVDepth	GeoMIM	Swin-B	512 \times 1408	✓	0.523	0.605	0.470	0.260	0.377	0.254	0.195

Table 2: Comparison on nuScenes val set. \dagger denotes our implementation with the official code.

Pretrain	3D-Segmentation		Waymo		nuImages	
	mIoU $^{\text{val}}$	mIoU $^{\text{test}}$	LET-3D APL	LET-3D AP	AP $^{\text{box}}$	AP $^{\text{mask}}$
	TPVFormer [22]		DfM [43]		Mask-RCNN [15]	
Supervised [33]	66.4	68.3	31.5	44.6	49.0	41.3
Self-supervised [30]	65.0	66.3	34.8	49.5	51.5	41.8
GeoMIM	68.9	70.5	37.8	52.5	52.9	44.4

Table 3: Transfer learning results on 3D-segmentation with TPVFormer (left), Open Waymo 3D detection with DfM (middle), and nuImages object detection and segmentation with Mask-RCNN (right).

Pretrain	NDS \uparrow	mAP \uparrow	mATE \downarrow	mAOE \downarrow
Supervised [33]	0.406	0.326	0.665	0.546
EsViT [25]	0.389	0.305	0.699	0.516
UniCL [48]	0.396	0.314	0.694	0.596
MixMAE [30]	0.443	0.374	0.647	0.418
GeoMIM	0.472	0.397	0.614	0.395

Table 4: Comparison with previous pretraining methods on nuScenes with BEVDet.

stream tasks.

For the 3D detection on the nuScenes dataset, we utilize the BEVDepth [28] framework with an input size of 512 \times 1408 for comparison with other state-of-the-art approaches. For ablation studies, we use the BEVDet [20] framework with an input size of 256 \times 704. For 3D detection on the Open Waymo dataset, the DfM [43, 42] framework is utilized. For 3D segmentation on the nuScenes dataset, we utilize the recent TPVFormer [22] for finetuning. We use Mask-RCNN [15] for object detection and instance segmentation on nuImages. We follow those frameworks’ default settings for finetuning, and include the detailed hyperparameters settings in the supplementary.

5. Main Results

In this section, we compare our GeoMIM to prior arts on various benchmarks. We first conduct comparisons between GeoMIM and previous pretraining approaches in Sec. 5.1. We then compare our best results with state-of-the-art results on the nuScenes 3D detection benchmark in Sec. 5.2. To show the transferability of GeoMIM, we present the transfer

learning results on other 3 benchmarks in Sec. 5.3. We finally show the quantitative results in Sec. 7.

5.1. Comparison with previous camera-based pre-training methods

We compare our pretraining method, GeoMIM, with previous pretraining approaches to demonstrate its effectiveness in multi-view camera-based 3D detection. Four pretraining approaches for camera-based models are utilized, including the supervised pretraining on ImageNet-1K [33], the contrastive approach EsViT [25], the multi-modal approach UniCL [48], and masked-image-modeling approach MixMAE [30]. Using the BEVDet framework with input size of 256 \times 704, we finetune the pretrained Swin-B [33] models on nuScenes [7] and compare their performances in Tab. 4. Our approach outperforms other compared approaches in terms of all reported metrics, demonstrating the effectiveness of our pretraining method.

Particularly, our approach achieves 0.472 NDS (NuScenes Detection Score), 2.9% better than the self-supervised pretraining. Notably, our approach is much better at predicting translation, demonstrating a 3.3% improvement in mATE, which shows that our geometry-enhanced pretraining can help more with localization. Surprisingly, while the contrastive or multi-modal approaches perform much better than the supervised ImageNet pretraining on various 2D visual tasks, they fail to improve the ImageNet-supervised pretraining on the 3D detection task.

Methods	Backbone	Image Size	Extra Data	TTA	mAP \uparrow	NDS \uparrow	mATE \downarrow	mASE \downarrow	mAOE \downarrow	mAVE \downarrow	mAAE \downarrow
FCOS3D [44]	R101-DCN	900 \times 1600	\times	\checkmark	0.358	0.428	0.690	0.249	0.452	1.434	0.124
DETR3D [45]	V2-99	900 \times 1600	\checkmark	\checkmark	0.412	0.479	0.641	0.255	0.394	0.845	0.133
UVTR [27]	V2-99	900 \times 1600	\checkmark	\times	0.472	0.551	0.577	0.253	0.391	0.508	0.123
BEVFormer [29]	V2-99	900 \times 1600	\checkmark	\times	0.481	0.569	0.582	0.256	0.375	0.378	0.126
BEVDet4D [19]	Swin-B	900 \times 1600	\times	\checkmark	0.451	0.569	0.511	0.241	0.386	0.301	0.121
PolarFormer [24]	V2-99	900 \times 1600	\checkmark	\times	0.493	0.572	0.556	0.256	0.364	0.439	0.127
PETrv2 [32]	GLOM-like	640 \times 1600	\times	\times	0.512	0.592	0.547	0.242	0.360	0.367	0.126
BEVDepth [28]	ConvNeXt-B	640 \times 1600	\times	\times	0.520	0.609	0.445	0.243	0.352	0.347	0.127
BEVStereo [26]	V2-99	640 \times 1600	\checkmark	\times	0.525	0.610	0.431	0.246	0.358	0.357	0.138
SOLOFusion [36]	ConvNeXt-B	640 \times 1600	\times	\times	0.540	0.619	0.453	0.257	0.376	0.276	0.148
BEVDistill [8]	ConvNeXt-B	640 \times 1600	\times	\times	0.498	0.594	0.472	0.247	0.378	0.326	0.125
GeoMIM	Swin-B	512 \times 1408	\times	\times	0.547	0.626	0.413	0.241	0.421	0.272	0.127
GeoMIM	Swin-L	512 \times 1408	\times	\times	0.561	0.644	0.400	0.238	0.348	0.255	0.120

Table 5: Comparison on nuScenes test set. “Extra data” denotes depth pretraining. “TTA” denotes test-time augmentation.

5.2. Comparison with state-of-the-art results

Tab. 2 shows the comparison of our approach with state-of-the-art methods on nuScenes val set. Our approach achieves state-of-the-art results of 0.605 NDS and 0.523 mAP, demonstrating substantial 2.3% NDS and 4.0% mAP improvements over SOLOFusion [36]. Particularly, the most improvement of NDS comes from the mATE, which improves SOLOFusion by 2.7%. Compared to BEVDepth [28] using the same Swin-B backbone, we improve the NDS and mAP by 5.0% and 5.7% respectively.

On the test set, our single model achieves 64.4% NDS and 56.1% mAP without using extra data and test-time augmentation, which are 2.5% and 2.1% better than the previous state-of-the-art results. Notably, this model performs best among all reported metrics. Compared to BEVStereo [26], the most significant improvement of NDS comes from the mAVE (10.2%), which shows that our geometry-enhanced pretraining is not only better for localization but also improves the velocity estimation. Compared to SOLOFusion, we largely improve the mATE metric (5.3%), showing that our pretraining is beneficial for localization.

We also show that our pretraining is scalable in terms of model size. In particular, on the test set, we obtain 1.8% NDS and 1.5% mAP gains by using the larger Swin-L [33] backbone.

5.3. Transfer to various 3D understanding tasks

In this section, we evaluate the transferability of our approach to other datasets and tasks with different frameworks. We use three benchmarks, 3D segmentation on nuScenes dataset [7], 3D detection on Open Waymo dataset [40], and object detection and instance segmentation on nuImages dataset.

As shown in Tab. 3, our approach achieves superior results on all three benchmarks, demonstrating the transferability of our pretraining method. Particularly, on the 3D segmentation task, our approach achieves 68.9% mIoU on the nuScenes val set, surpassing the ImageNet-supervised pretraining [33] results by a large margin. Note that, unlike the 3D detection task, the self-supervised pretraining [30] fails to

improve the supervised pretraining because the segmentation task relies more on semantic understanding. In comparison, GeoMIM improves the ImageNet-supervised pretraining for 2.5% mIoU. On the nuScenes test test, we achieve state-of-the-art results, 1.1% mIoU better than the previous best camera-based results in TPVFormer [22]. Moreover, our pretrained backbone can also transfer to datasets that differ from that used in pretraining. On Open Waymo detection benchmark, our pretraining improves the MixMAE [30] self-supervised pretrained model by 3.0%/3.0% on LET-3D APL/AP.

Apart from the 3D perception task, we show that our pretrain can also transfer to 2D object detection and instance segmentation tasks. As shown in Tab. 3 (right), GeoMIM improves the self-supervised pretraining by 1.4% AP^{box} and 2.6% AP^{mask}.

6. Ablation Studies

In this section, we conduct ablation studies to evaluate the impact of different design choices on the performance of our proposed GeoMIM on the multi-view camera-based 3D detection task. Unless otherwise specified, we use the Swin-B [33] backbone and pretrain it for 6 epochs. We utilize the BEVDet [20] framework for finetuning the pretrained backbone and report the performance on the nuScenes val set [7]. The gray column indicates the final choice of GeoMIM.

Pretraining epochs and pretraining data. We explore the effect of pretraining epochs and pretraining data on GeoMIM. As shown in Tab. 6, we find that we can improve the mATE performance but degenerate the mAOE performance through 6 epochs of pretraining. Interestingly, if we pretrain for more epochs, mATE performance saturates but mAOE can be largely improved. Additionally, as shown in Tab. 7, the performance of all metrics gradually increases as we use more data for pretraining.

Mask ratio. We examine the effect of the mask ratio used in the masked image modeling training process on the performance. As shown in Tab. 8, we find that using a mask ratio of 0.5 performs best as a very high mask ratio causes

# epochs	NDS↑	mAP↑	mATE↓	mAOE↓
0	0.443	0.374	0.647	0.418
6	0.460	0.381	0.613	0.449
50	0.472	0.398	0.614	0.395
100	0.475	0.400	0.615	0.390

Table 6: Ablation of pretraining epochs.

% data	NDS↑	mAP↑	mATE↓	mAOE↓
0	0.443	0.374	0.647	0.418
10	0.426	0.356	0.646	0.519
50	0.445	0.367	0.623	0.465
100	0.460	0.381	0.613	0.449

Table 7: Ablation of the percentage of the pretraining data.

Mask ratio	NDS↑	mAP↑	mATE↓	mAOE↓
0.25	0.454	0.375	0.622	0.447
0.5	0.460	0.381	0.613	0.449
0.75	0.453	0.375	0.642	0.452

Table 8: Ablation of the mask ratio.

the pretext task too hard while a low mask ratio causes the pretext task too easy.

Pretraining with distillation or other reconstruction targets. We compare the performance of GeoMIM with different learning targets, including RGB pixels [14] and the voxelized LiDAR points. Moreover, we use the depth ground truth derived from the LiDAR point cloud for depth pretraining [35]. Following previous works, we also use the LiDAR BEV features for conducting distillation pretraining [13, 8, 17]. We initialize the backbone with MixMAE [30] self-supervised pretraining and use its results for comparison. All the pretraining experiments are conducted on the nuScenes dataset.

As shown in Tab. 9, we find the depth or distillation pretraining fails to improve the MixMAE results. Those two pretraining methods are beneficial for object localization to improve the mATE metric, but degenerate mAOE a lot. Using the RGB pixels as the reconstruction targets like MAE is also unable to improve the NDS. The main reason is that the nuScenes dataset is much less diverse than the widely used ImageNet-1K [10] dataset, and as a result, the model is easy to overfit the training data. Moreover, though the LiDAR points contain rich geometry information, we find that directly using the voxelized LiDAR points as the reconstruction targets also fails to improve the MixMAE results. As stated in Sec. 1, the LiDAR voxels are sparse and noisy. Using them as the reconstruction targets results in unstable pretraining. In comparison, we use a pretrained LiDAR detection model to extract more meaningful BEV features as the reconstruction targets, which can not only transfer the rich geometry information to the camera-based model but also avoid the noise problem of directly reconstructing LiDAR voxels.

Ablation on the decode designs. We investigate the impact of the proposed decoder designs on the final performance. Apart from reporting results on the nuScenes dataset,

Approach	NDS↑	mAP↑	mATE↓	mAOE↓
SSL Pretrain [30]	0.443	0.374	0.647	0.418
+ Depth Pretrain	0.426	0.350	0.623	0.535
+ Distill Pretrain	0.434	0.360	0.626	0.533
+ LiDAR reconstruction	0.439	0.352	0.637	0.483
+ RGB reconstruction	0.440	0.376	0.632	0.484
+ GeoMIM	0.460	0.381	0.613	0.449

Table 9: Ablation of pretraining setups and targets.

Decouple	CVA	Cam.	nuScenes		Waymo	
			NDS	mAP	APL	AP
✓	✓	✓	46.0	38.1	37.0	51.6
✗	✓	✓	45.2	37.8	36.4	51.4
✓	✗	✓	45.5	38.1	36.8	51.3
✓	✗	✗	45.1	37.9	36.0	51.0

Table 10: Ablation of the decoder designs. “Cam.” denotes the camera-aware design.

Backbone	Parameters	NDS↑	mAP↑	mATE↓	mAOE↓
Swin-B	88M	0.460	0.381	0.613	0.449
Swin-L	197M	0.478	0.398	0.609	0.371

Table 11: Ablation of the model sizes.

we also report the performance on the Open Waymo dataset to show how the design choices affect the transferability of the pretraining.

As shown in Tab. 10, we find that using a decoupled decoder to separately decode the depth and semantic features can improve the NDS for 0.8%, compared to using one decoder to jointly decode them. The decoupled branches force the geometry branch to focus on inferring depth from geometry cues, which maximizes the utilization of the model capacity. Moreover, removing the Cross-View Attention (CVA) blocks results in a 0.5% performance drop because of lacking cross-view interaction for better BEV feature inference. Additionally, we find that further removing the camera-aware design leads to 0.8% LET-3D APL drop on the Open Waymo dataset. As the pretraining and finetuning might be conducted on different datasets, using camera-aware depth reconstruction is beneficial for the transferability of the pretraining.

Ablation on backbone size. We investigate the scalability of our pretraining method and show the results in Tab. 11. Our GeoMIM is capable to be scaled up to Swin-L model with 200M parameters. The pretrained Swin-L [33] improves Swin-B on all reported metrics, especially the mAOE performance (7.8%).

7. Qualitative Evaluation

Reconstruction visualization. We show the decoded geometry features and semantic features during pretraining in Fig. 3. Although a high mask ratio is utilized, meaningful patterns can still be observed in the decoded features. In particular, we find the geometry branch focuses on the structure of the driving scenes while the semantic branch can be activated by different semantic regions, including roads, cars,

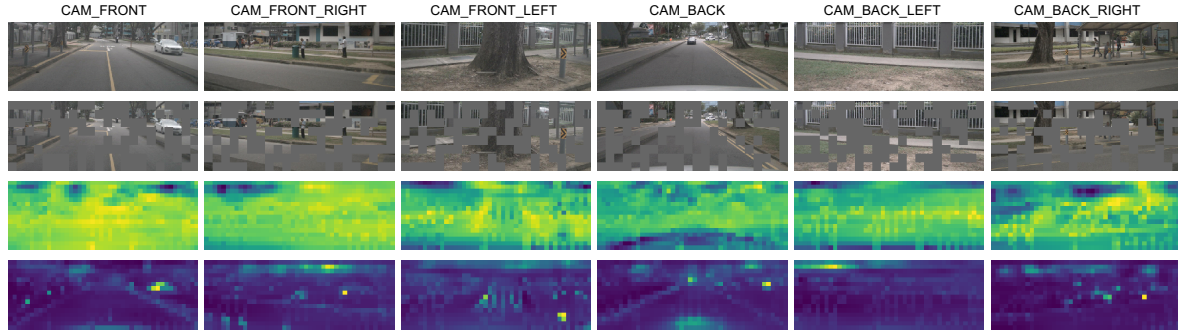


Figure 3: Example results on nuScenes val images. From top to bottom, the rows are the camera-view image, masked camera-view image, decoded semantic features, and decoded geometry features.




Figure 4: Reconstruction results on nuScenes val images. For each triplet, we show the reconstructed BEV features (left), LiDAR BEV features, and LiDAR point cloud (right) in BEV.

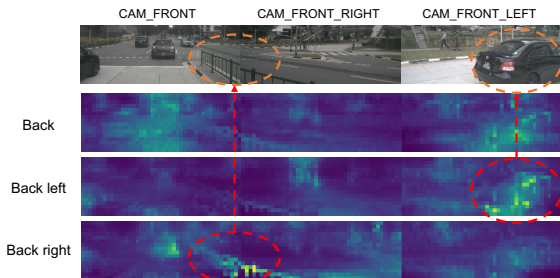


Figure 5: CVA blocks' attention maps across different cameras. Our CVA enables the cameras at back to attend to the semantic parts in front views.

trees, etc. Furthermore, we visualize the reconstructed BEV features in Fig. 4. The reconstructed BEV features can well restore the LiDAR model's BEV features, including the road structures and the semantic features.

Cross-view attention maps. We visualize the attention maps of the Cross-View Attention (CVA) blocks in Fig. 5. Through the cross-view interaction, one view is able to attend to the semantic parts of other views.

Convergence curve. We show the convergence curve of different pretraining in Fig. 6. Our pretraining can largely improve the self-supervised [30] and supervised [33] pretraining in terms of the convergence speed, which can match the self-supervised pretraining's final results with only half of the iterations.

8. Conclusion

In this paper, we proposed a pretraining method, GeoMIM, for multi-view camera-based 3D detection. By leveraging the knowledge of a pretrained LiDAR model in a




Figure 6: Performance curves of different pretraining methods. “SL” and “SSL” denote the ImageNet-supervised and MixMAE self-supervised pretraining respectively.

pretrain-finetune paradigm, GeoMIM aims to transfer its rich geometry knowledge to the camera-based model. Specifically, GeoMIM reconstructs BEV features from masked images via a novel decoder and a cross-view attention mechanism. We demonstrate that GeoMIM significantly outperforms existing state-of-the-art methods on the nuScenes dataset, achieving state-of-the-art results in both camera-based 3D detection and segmentation tasks. Moreover, we verify that the pretrained model can be transferred to the Waymo Open dataset, further showing its effectiveness and generality.

Limitations Despite the promising results, GeoMIM also has some limitations. First, GeoMIM requires a large amount of labeled data for pretraining, which may not be available in some applications. Second, GeoMIM relies on the quality of the LiDAR model's BEV features, which may not always be accurate or complete. Overall, while GeoMIM shows great potential, further research is needed to address these limitations and improve its applicability in a wider range of applications.

References

- [1] Mahmoud Assran, Mathilde Caron, Ishan Misra, Piotr Bojanowski, Florian Bordes, Pascal Vincent, Armand Joulin, Mike Rabbat, and Nicolas Ballas. Masked siamese networks for label-efficient learning. In *ECCV*, 2022. 2
- [2] Sara Atito, Muhammad Awais, and Josef Kittler. Sit: Self-supervised vision transformer. *arXiv:2104.03602*, 2021. 2
- [3] Roman Bachmann, David Mizrahi, Andrei Atanov, and Amir Zamir. Multimaec: Multi-modal multi-task masked autoencoders. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXVII*, pages 348–367. Springer, 2022. 2
- [4] Alexei Baevski, Wei-Ning Hsu, Qiantong Xu, Arun Babu, Jiatao Gu, and Michael Auli. Data2vec: A general framework for self-supervised learning in speech, vision and language. *arXiv:2202.03555*, 2022. 2
- [5] Xuyang Bai, Zeyu Hu, Xinge Zhu, Qingqiu Huang, Yilun Chen, Hongbo Fu, and Chiew-Lan Tai. Transfusion: Robust lidar-camera fusion for 3d object detection with transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1090–1099, 2022. 4
- [6] Hangbo Bao, Li Dong, and Furu Wei. Beit: Bert pre-training of image transformers. In *ICLR*, 2021. 2
- [7] Holger Caesar, Varun Bankiti, Alex H Lang, Sourabh Vora, Venice Erin Liong, Qiang Xu, Anush Krishnan, Yu Pan, Giancarlo Baldan, and Oscar Beijbom. nuscenes: A multi-modal dataset for autonomous driving. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 11621–11631, 2020. 2, 4, 5, 6
- [8] Zehui Chen, Zhenyu Li, Shiquan Zhang, Liangji Fang, Qinhong Jiang, and Feng Zhao. Bevdistill: Cross-modal bev distillation for multi-view 3d object detection. *arXiv preprint arXiv:2211.09386*, 2022. 1, 2, 6, 7
- [9] Zhiyu Chong, Xinzhu Ma, Hong Zhang, Yuxin Yue, Haojie Li, Zhihui Wang, and Wanli Ouyang. Monodistill: Learning spatial features for monocular 3d object detection. *arXiv preprint arXiv:2201.10830*, 2022. 1, 2
- [10] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, K. Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *CVPR*, 2009. 4, 7
- [11] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019. 2
- [12] Xiaoyi Dong, Jianmin Bao, Ting Zhang, Dongdong Chen, Weiming Zhang, Lu Yuan, Dong Chen, Fang Wen, and Nenghai Yu. Peco: Perceptual codebook for bert pre-training of vision transformers. *arXiv:2111.12710*, 2021. 2
- [13] Xiaoyang Guo, Shaoshuai Shi, Xiaogang Wang, and Hongsheng Li. Liga-stereo: Learning lidar geometry aware representations for stereo-based 3d detector. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 3153–3163, 2021. 1, 2, 7
- [14] Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross B. Girshick. Masked autoencoders are scalable vision learners. In *CVPR*, 2021. 1, 2, 4, 7
- [15] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn. In *Proceedings of the IEEE international conference on computer vision*, pages 2961–2969, 2017. 5
- [16] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015. 2
- [17] Yu Hong, Hang Dai, and Yong Ding. Cross-modality knowledge distillation network for monocular 3d object detection. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part X*, pages 87–104. Springer, 2022. 1, 2, 7
- [18] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 7132–7141, 2018. 4
- [19] Junjie Huang and Guan Huang. Bevdet4d: Exploit temporal cues in multi-camera 3d object detection. *arXiv preprint arXiv:2203.17054*, 2022. 1, 5, 6
- [20] Junjie Huang, Guan Huang, Zheng Zhu, and Dalong Du. Bevdet: High-performance multi-camera 3d object detection in bird-eye-view. *arXiv preprint arXiv:2112.11790*, 2021. 1, 2, 4, 5, 6
- [21] Peixiang Huang, Li Liu, Renrui Zhang, Song Zhang, Xinli Xu, Baichao Wang, and Guoyi Liu. Tig-bev: Multi-view bev 3d object detection via target inner-geometry learning. *arXiv preprint arXiv:2212.13979*, 2022. 1
- [22] Yuanhui Huang, Wenzhao Zheng, Yunpeng Zhang, Jie Zhou, and Jiwen Lu. Tri-perspective view for vision-based 3d semantic occupancy prediction. *arXiv preprint arXiv:2302.07817*, 2023. 2, 5, 6
- [23] Wei-Chih Hung, Henrik Kretschmar, Vincent Casser, Jyh-Jing Hwang, and Dragomir Anguelov. Let-3d-ap: Longitudinal error tolerant 3d average precision for camera-only 3d detection. *arXiv preprint arXiv:2206.07705*, 2022. 4
- [24] Yanqin Jiang, Li Zhang, Zhenwei Miao, Xiatian Zhu, Jin Gao, Weiming Hu, and Yu-Gang Jiang. Polarformer: Multi-camera 3d object detection with polar transformers. *arXiv preprint arXiv:2206.15398*, 2022. 5, 6
- [25] Chunyuan Li, Jianwei Yang, Pengchuan Zhang, Mei Gao, Bin Xiao, Xiyang Dai, Lu Yuan, and Jianfeng Gao. Efficient self-supervised vision transformers for representation learning. In *ICLR*, 2022. 5
- [26] Yinhao Li, Han Bao, Zheng Ge, Jinrong Yang, Jianjian Sun, and Zeming Li. Bevestereo: Enhancing depth estimation in multi-view 3d object detection with dynamic temporal stereo. *arXiv preprint arXiv:2209.10248*, 2022. 1, 2, 5, 6
- [27] Yanwei Li, Yilun Chen, Xiaojuan Qi, Zeming Li, Jian Sun, and Jiaya Jia. Unifying voxel-based representation with transformer for 3d object detection. *arXiv preprint arXiv:2206.00630*, 2022. 1, 2, 5, 6
- [28] Yinhao Li, Zheng Ge, Guanyi Yu, Jinrong Yang, Zengran Wang, Yukang Shi, Jianjian Sun, and Zeming Li. Bevedepth: Acquisition of reliable depth for multi-view 3d object detection. *arXiv preprint arXiv:2206.10092*, 2022. 1, 2, 4, 5, 6
- [29] Zhiqi Li, Wenhao Wang, Hongyang Li, Enze Xie, Chonghao Sima, Tong Lu, Yu Qiao, and Jifeng Dai. Bevformer: Learning bird’s-eye-view representation from multi-camera images via spatiotemporal transformers. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part IX*, pages 1–18. Springer, 2022. 1, 2, 5, 6

- [30] Jihao Liu, Xin Huang, Yu Liu, and Hongsheng Li. Mixmim: Mixed and masked image modeling for efficient visual representation learning. *arXiv preprint arXiv:2205.13137*, 2022. 1, 2, 4, 5, 6, 7, 8
- [31] Yingfei Liu, Tiancai Wang, Xiangyu Zhang, and Jian Sun. Petr: Position embedding transformation for multi-view 3d object detection. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXVII*, pages 531–548. Springer, 2022. 1, 5
- [32] Yingfei Liu, Junjie Yan, Fan Jia, Shuailin Li, Qi Gao, Tiancai Wang, Xiangyu Zhang, and Jian Sun. Petr2: A unified framework for 3d perception from multi-camera images. *arXiv preprint arXiv:2206.01256*, 2022. 1, 5, 6
- [33] Ze Liu, Yutong Lin, Yue Cao, Han Hu, Yixuan Wei, Zheng Zhang, Stephen Lin, and Baining Guo. Swin transformer: Hierarchical vision transformer using shifted windows. In *ICCV*, 2021. 1, 2, 4, 5, 6, 7, 8
- [34] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *ICLR*, 2017. 4
- [35] Dennis Park, Rares Ambrus, Vitor Guizilini, Jie Li, and Adrien Gaidon. Is pseudo-lidar needed for monocular 3d object detection? In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 3142–3152, 2021. 7
- [36] Jinhyung Park, Chenfeng Xu, Shijia Yang, Kurt Keutzer, Kris Kitani, Masayoshi Tomizuka, and Wei Zhan. Time will tell: New outlooks and a baseline for temporal multi-view 3d object detection. *arXiv preprint arXiv:2210.02443*, 2022. 1, 2, 5, 6
- [37] Jonah Philion and Sanja Fidler. Lift, splat, shoot: Encoding images from arbitrary camera rigs by implicitly unprojecting to 3d. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XIV 16*, pages 194–210. Springer, 2020. 2, 3
- [38] Aditya Ramesh, Mikhail Pavlov, Gabriel Goh, Scott Gray, Chelsea Voss, Alec Radford, Mark Chen, and Ilya Sutskever. Zero-shot text-to-image generation. In *ICML*, 2021. 2
- [39] Jason Tyler Rolfe. Discrete variational autoencoders. In *ICLR*, 2016. 2
- [40] Pei Sun, Henrik Kretschmar, Xerxes Dotiwalla, Aurelien Chouard, Vijaysai Patnaik, Paul Tsui, James Guo, Yin Zhou, Yuning Chai, Benjamin Caine, et al. Scalability in perception for autonomous driving: Waymo open dataset. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 2446–2454, 2020. 2, 4, 6
- [41] Ashish Vaswani, Noam M. Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *NeurIPS*, 2017. 3
- [42] Tai Wang, Qing Lian, Chenming Zhu, Xinge Zhu, and Wenwei Zhang. Mv-fcos3d++: Multi-view camera-only 4d object detection with pretrained monocular backbones. *arXiv preprint arXiv:2207.12716*, 2022. 5
- [43] Tai Wang, Jiangmiao Pang, and Dahua Lin. Monocular 3d object detection with depth from motion. In *Computer Vision–ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part IX*, pages 386–403. Springer, 2022. 5
- [44] Tai Wang, Xinge Zhu, Jiangmiao Pang, and Dahua Lin. Fcos3d: Fully convolutional one-stage monocular 3d object detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 913–922, 2021. 1, 2, 6
- [45] Yue Wang, Vitor Campagnolo Guizilini, Tianyuan Zhang, Yilun Wang, Hang Zhao, and Justin Solomon. Detr3d: 3d object detection from multi-view images via 3d-to-2d queries. In *Conference on Robot Learning*, pages 180–191. PMLR, 2022. 1, 2, 5, 6
- [46] Tete Xiao, Ilija Radosavovic, Trevor Darrell, and Jitendra Malik. Masked visual pre-training for motor control. *arXiv:2111.11429*, 2022. 2
- [47] Zhenda Xie, Zheng Zhang, Yue Cao, Yutong Lin, Jianmin Bao, Zhuliang Yao, Qi Dai, and Han Hu. Simmim: A simple framework for masked image modeling. In *CVPR*, 2021. 2, 4
- [48] Jianwei Yang, Chunyuan Li, Pengchuan Zhang, Bin Xiao, Ce Liu, Lu Yuan, and Jianfeng Gao. Unified contrastive learning in image-text-label space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 19163–19173, 2022. 5
- [49] Renrui Zhang, Ziyu Guo, Peng Gao, Rongyao Fang, Bin Zhao, Dong Wang, Yu Qiao, and Hongsheng Li. Point-m2ae: multi-scale masked autoencoders for hierarchical point cloud pre-training. *arXiv preprint arXiv:2205.14401*, 2022. 2
- [50] Renrui Zhang, Han Qiu, Tai Wang, Xuanzhuo Xu, Ziyu Guo, Yu Qiao, Peng Gao, and Hongsheng Li. Monodetr: Depth-aware transformer for monocular 3d object detection. *arXiv preprint arXiv:2203.13310*, 2022. 2
- [51] Jinghao Zhou, Chen Wei, Huiyu Wang, Wei Shen, Cihang Xie, Alan Yuille, and Tao Kong. ibot: Image bert pre-training with online tokenizer. In *ICLR*, 2021. 2